

Universidad ICESI

Facultad de ingeniería y diseño

Maestría en ciencia de datos

Proyecto final - Fundamentos de analítica

Laura Loaiza

```
!pip install lazy
!pip install lazypredict
!pip install hyperopt
!pip install pyclustering

Collecting lazy
  Downloading lazy-1.5-py2.py3-none-any.whl (5.0 kB)
Installing collected packages: lazy
Successfully installed lazy-1.5
Collecting lazypredict
  Downloading lazypredict-0.2.12-py2.py3-none-any.whl (12 kB)
Requirement already satisfied: pandas in c:\users\mq200-0489\
anaconda3\lib\site-packages (from lazypredict) (1.4.4)
Requirement already satisfied: scikit-learn in c:\users\mq200-0489\
anaconda3\lib\site-packages (from lazypredict) (1.0.2)
Requirement already satisfied: joblib in c:\users\mq200-0489\
anaconda3\lib\site-packages (from lazypredict) (1.1.0)
Collecting xgboost
  Downloading xgboost-1.7.5-py3-none-win_amd64.whl (70.9 MB)
----- 70.9/70.9 MB 376.6 kB/s
eta 0:00:00
Requirement already satisfied: tqdm in c:\users\mq200-0489\anaconda3\
lib\site-packages (from lazypredict) (4.64.1)
Requirement already satisfied: click in c:\users\mq200-0489\
anaconda3\lib\site-packages (from lazypredict) (8.0.4)
Collecting lightgbm
  Downloading lightgbm-3.3.5-py3-none-win_amd64.whl (1.0 MB)
----- 1.0/1.0 MB 968.3 kB/s
eta 0:00:00
Requirement already satisfied: colorama in c:\users\mq200-0489\
anaconda3\lib\site-packages (from click->lazypredict) (0.4.5)
Requirement already satisfied: numpy in c:\users\mq200-0489\
anaconda3\lib\site-packages (from lightgbm->lazypredict) (1.21.5)
Requirement already satisfied: scipy in c:\users\mq200-0489\
anaconda3\lib\site-packages (from lightgbm->lazypredict) (1.9.1)
Requirement already satisfied: wheel in c:\users\mq200-0489\
anaconda3\lib\site-packages (from lightgbm->lazypredict) (0.37.1)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\
mq200-0489\anaconda3\lib\site-packages (from scikit-learn-
```

```

>lazypredict) (2.2.0)
Requirement already satisfied: python-dateutil>=2.8.1 in c:\users\
mqa200-0489\anaconda3\lib\site-packages (from pandas->lazypredict)
(2.8.2)
Requirement already satisfied: pytz>=2020.1 in c:\users\mqa200-0489\
anaconda3\lib\site-packages (from pandas->lazypredict) (2022.1)
Requirement already satisfied: six>=1.5 in c:\users\mqa200-0489\
anaconda3\lib\site-packages (from python-dateutil>=2.8.1->pandas-
>lazypredict) (1.16.0)
Installing collected packages: xgboost, lightgbm, lazypredict
Successfully installed lazypredict-0.2.12 lightgbm-3.3.5 xgboost-1.7.5
Collecting hyperopt
  Downloading hyperopt-0.2.7-py2.py3-none-any.whl (1.6 MB)
----- 1.6/1.6 MB 498.5 kB/s
eta 0:00:00
Requirement already satisfied: numpy in c:\users\mqa200-0489\
anaconda3\lib\site-packages (from hyperopt) (1.21.5)
Requirement already satisfied: scipy in c:\users\mqa200-0489\
anaconda3\lib\site-packages (from hyperopt) (1.9.1)
Requirement already satisfied: six in c:\users\mqa200-0489\anaconda3\
lib\site-packages (from hyperopt) (1.16.0)
Requirement already satisfied: networkx>=2.2 in c:\users\mqa200-0489\
anaconda3\lib\site-packages (from hyperopt) (2.8.4)
Requirement already satisfied: future in c:\users\mqa200-0489\
anaconda3\lib\site-packages (from hyperopt) (0.18.2)
Collecting py4j
  Downloading py4j-0.10.9.7-py2.py3-none-any.whl (200 kB)
----- 200.5/200.5 kB 347.8 kB/s
eta 0:00:00
Requirement already satisfied: cloudpickle in c:\users\mqa200-0489\
anaconda3\lib\site-packages (from hyperopt) (2.0.0)
Requirement already satisfied: tqdm in c:\users\mqa200-0489\anaconda3\
lib\site-packages (from hyperopt) (4.64.1)
Requirement already satisfied: colorama in c:\users\mqa200-0489\
anaconda3\lib\site-packages (from tqdm->hyperopt) (0.4.5)
Installing collected packages: py4j, hyperopt
Successfully installed hyperopt-0.2.7 py4j-0.10.9.7
Collecting pylustering
  Downloading pylustering-0.10.1.2.tar.gz (2.6 MB)
----- 2.6/2.6 MB 218.0 kB/s
eta 0:00:00
  Preparing metadata (setup.py): started
  Preparing metadata (setup.py): finished with status 'done'
Requirement already satisfied: scipy>=1.1.0 in c:\users\mqa200-0489\
anaconda3\lib\site-packages (from pylustering) (1.9.1)
Requirement already satisfied: matplotlib>=3.0.0 in c:\users\mqa200-
0489\anaconda3\lib\site-packages (from pylustering) (3.5.2)
Requirement already satisfied: numpy>=1.15.2 in c:\users\mqa200-0489\
anaconda3\lib\site-packages (from pylustering) (1.21.5)

```

```
Requirement already satisfied: Pillow>=5.2.0 in c:\users\mqa200-0489\anaconda3\lib\site-packages (from pyclustering) (9.2.0)
Requirement already satisfied: packaging>=20.0 in c:\users\mqa200-0489\anaconda3\lib\site-packages (from matplotlib>=3.0.0->pyclustering) (21.3)
Requirement already satisfied: fonttools>=4.22.0 in c:\users\mqa200-0489\anaconda3\lib\site-packages (from matplotlib>=3.0.0->pyclustering) (4.25.0)
Requirement already satisfied: cycler>=0.10 in c:\users\mqa200-0489\anaconda3\lib\site-packages (from matplotlib>=3.0.0->pyclustering) (0.11.0)
Requirement already satisfied: kiwisolver>=1.0.1 in c:\users\mqa200-0489\anaconda3\lib\site-packages (from matplotlib>=3.0.0->pyclustering) (1.4.2)
Requirement already satisfied: pyparsing>=2.2.1 in c:\users\mqa200-0489\anaconda3\lib\site-packages (from matplotlib>=3.0.0->pyclustering) (3.0.9)
Requirement already satisfied: python-dateutil>=2.7 in c:\users\mqa200-0489\anaconda3\lib\site-packages (from matplotlib>=3.0.0->pyclustering) (2.8.2)
Requirement already satisfied: six>=1.5 in c:\users\mqa200-0489\anaconda3\lib\site-packages (from python-dateutil>=2.7->matplotlib>=3.0.0->pyclustering) (1.16.0)
Building wheels for collected packages: pyclustering
  Building wheel for pyclustering (setup.py): started
  Building wheel for pyclustering (setup.py): finished with status 'done'
  Created wheel for pyclustering: filename=pyclustering-0.10.1.2-py3-none-any.whl size=2395106
  sha256=6cc3f22cdb19dd5b73c6178c1c7b93b960beaf92f416ef1fa2ac83d694825898
  Stored in directory: c:\users\mqa200-0489\appdata\local\pip\cache\wheels\e0\56\c2\abb6866a3fcd8a55862f1df8a18f57805c3a78fed9a9023cb9
Successfully built pyclustering
Installing collected packages: pyclustering
Successfully installed pyclustering-0.10.1.2
```

```
import pandas as pd
import numpy as np
import warnings
warnings.simplefilter("ignore")
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import pandas as pd
import math
from collections import Counter
from sklearn.cluster import KMeans
from sklearn.metrics import confusion_matrix, accuracy_score, silhouette_samples, silhouette_score, calinski_harabasz_score
```

```

from sklearn import preprocessing
from sklearn.decomposition import PCA
from dataprep.eda import create_report
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn.metrics import cohen_kappa_score
from lazypredict.Supervised import LazyClassifier
import xgboost as xgb
from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from hyperopt import hp, tpe, fmin
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.datasets import make_classification
from sklearn.cluster import KMeans
from sklearn.metrics import calinski_harabasz_score
from sklearn.metrics import silhouette_samples
import matplotlib.pyplot as plt
import lightgbm as lgb
from sklearn.cluster import AgglomerativeClustering
from scipy.cluster.hierarchy import linkage
from scipy.cluster.hierarchy import dendrogram

""

""

#Cargamos el dataset
data=pd.read_csv("imports-85.data", header=None)

#Renombramos las columnas con nombres mas apropiados para los
atributos
data.columns = [
    "symboling", "normalized-losses", "make", "fuel-type",
    "aspiration", "num-of-doors", "body-style", "drive-wheels",
    "engine-location", "wheel-base", "length", "width", "height",
    "curb-weight", "engine-type", "num-of-cylinders",
    "engine-size", "fuel-system", "bore", "stroke", "compression-
    ratio", "horsepower", "peak-rpm", "city-mpg", "highway-mpg",
    "price"
]

#reemplazamos los signos de interrogación (?) con nan
data=data.replace("?",np.nan)

data.head()

```

symboling	normalized-losses	make	fuel-type	aspiration	num-	
of-doors \						
0	3	NaN	alfa-romero	gas	std	
two						
1	3	NaN	alfa-romero	gas	std	
two						
2	1	NaN	alfa-romero	gas	std	
two						
3	2	164.00	audi	gas	std	
four						
4	2	164.00	audi	gas	std	
four						
body-style	drive-wheels	engine-location	wheel-base	...	engine-	
size \						
0	convertible	rwd	front	88.60	...	
130						
1	convertible	rwd	front	88.60	...	
130						
2	hatchback	rwd	front	94.50	...	
152						
3	sedan	fwd	front	99.80	...	
109						
4	sedan	4wd	front	99.40	...	
136						
fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm	
city-mpg \						
0	mpfi	3.47	2.68	9.00	111.00	5000.00
21						
1	mpfi	3.47	2.68	9.00	111.00	5000.00
21						
2	mpfi	2.68	3.47	9.00	154.00	5000.00
19						
3	mpfi	3.19	3.40	10.00	102.00	5500.00
24						
4	mpfi	3.19	3.40	8.00	115.00	5500.00
18						
highway-mpg	price					
0	27	13495.00				
1	27	16500.00				
2	26	16500.00				
3	30	13950.00				
4	22	17450.00				
[5 rows x 26 columns]						

Atributos:

1. symboling: -3, -2, -1, 0, 1, 2, 3.
2. normalized-losses: continua de 65 a 256.
3. make: alfa-romero, audi, bmw, chevrolet, dodge, honda, isuzu, jaguar, mazda, mercedes-benz, mercury, mitsubishi, nissan, peugot, plymouth, porsche, renault, saab, subaru, toyota, volkswagen, volvo
4. fuel-type: diesel, gas.
5. aspiration: std, turbo.
6. num-of-doors: four, two.
7. body-style: hardtop, wagon, sedan, hatchback, convertible.
8. drive-wheels: 4wd, fwd, rwd.
9. engine-location: front, rear.
10. wheel-base: continua de 86.6 a 120.9.
11. length: continua de 141.1 a 208.1.
12. width: continua de 60.3 a 72.3.
13. height: continua de 47.8 a 59.8.
14. curb-weight: continua de 1488 a 4066.
15. engine-type: dohc, dohcvt, l, ohc, ohcf, ohcv, rotor.
16. num-of-cylinders: eight, five, four, six, three, twelve, two.
17. engine-size: continua de 61 a 326.
18. fuel-system: 1bbl, 2bbl, 4bbl, idi, mfi, mpfi, spdi, spfi.
19. bore: continua de 2.54 a 3.94.
20. stroke: continua de 2.07 a 4.17.
21. compression-ratio: continua de 7 a 23.
22. horsepower: continua de 48 a 288.
23. peak-rpm: continua de 4150 a 6600.
24. city-mpg: continua de 13 a 49.

#verificamos que Dtype tiene cada atributo
data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   symboling              205 non-null    int64
1   normalized-losses      164 non-null    object
2   make                   205 non-null    object
3   fuel-type              205 non-null    object
4   aspiration              205 non-null    object
5   num-of-doors           203 non-null    object
6   body-style             205 non-null    object
7   drive-wheels           205 non-null    object
8   engine-location        205 non-null    object
9   wheel-base             205 non-null    float64
10  length                 205 non-null    float64
```

```

11 width                205 non-null    float64
12 height               205 non-null    float64
13 curb-weight          205 non-null    int64
14 engine-type          205 non-null    object
15 num-of-cylinders     205 non-null    object
16 engine-size          205 non-null    int64
17 fuel-system          205 non-null    object
18 bore                 201 non-null    object
19 stroke               201 non-null    object
20 compression-ratio    205 non-null    float64
21 horsepower           203 non-null    object
22 peak-rpm             203 non-null    object
23 city-mpg             205 non-null    int64
24 highway-mpg          205 non-null    int64
25 price                201 non-null    object
dtypes: float64(5), int64(5), object(16)
memory usage: 41.8+ KB

```

#Corregimos el formato de los atributos

```

data["normalized-losses"]=data["normalized-losses"].astype(float)
data["price"]=data["price"].astype(float)
data["bore"]=data["bore"].astype(float)
data["stroke"]=data["stroke"].astype(float)
data["horsepower"]=data["horsepower"].astype(float)
data["peak-rpm"]=data["peak-rpm"].astype(float)
data["price"]=data["price"].astype(float)
data["symboling"]=data["symboling"].astype(int)

```

#Revisamos los valores unicos de la variable dependiente Y

```

data["symboling"].unique()

array([ 3,  1,  2,  0, -1, -2])

```

#Análisis exploratorio de datos

```

create_report(data)

{"model_id":"","version_major":2,"version_minor":0}

```

#Matrx de correlaciones Spearman

```

correlation_matrix = data.corr(method='spearman')
correlation_matrix

```

	normalized-losses	wheel-base	length	width
height \				
normalized-losses	1.00	-0.11	0.02	0.11
0.39				
wheel-base	-0.11	1.00	0.91	0.81
0.63				
length	0.02	0.91	1.00	0.89

0.53					
width	0.11	0.81	0.89	1.00	
0.35					
height	-0.39	0.63	0.53	0.35	
1.00					
curb-weight	0.09	0.77	0.89	0.86	
0.35					
engine-size	0.08	0.65	0.78	0.77	
0.20					
bore	-0.06	0.54	0.64	0.61	
0.23					
stroke	0.09	0.22	0.18	0.24	-
0.03					
compression-ratio	-0.05	-0.13	-0.19	-0.15	
0.00					
horsepower	0.24	0.50	0.66	0.69	
0.01					
peak-rpm	0.30	-0.32	-0.27	-0.20	-
0.30					
city-mpg	-0.25	-0.49	-0.67	-0.69	-
0.07					
highway-mpg	-0.20	-0.54	-0.70	-0.70	-
0.13					
price	0.19	0.68	0.81	0.81	
0.26					

	curb-weight	engine-size	bore	stroke
compression-ratio	\			
normalized-losses	0.09	0.08	-0.06	0.09
-0.05				
wheel-base	0.77	0.65	0.54	0.22
-0.13				
length	0.89	0.78	0.64	0.18
-0.19				
width	0.86	0.77	0.61	0.24
-0.15				
height	0.35	0.20	0.23	-0.03
0.00				
curb-weight	1.00	0.88	0.70	0.16
-0.22				
engine-size	0.88	1.00	0.72	0.29
-0.23				
bore	0.70	0.72	1.00	-0.08
-0.17				
stroke	0.16	0.29	-0.08	1.00
-0.07				
compression-ratio	-0.22	-0.23	-0.17	-0.07
1.00				
horsepower	0.81	0.82	0.65	0.14

-0.36					
peak-rpm	-0.24	-0.28	-0.31	-0.07	
-0.03					
city-mpg	-0.81	-0.73	-0.62	-0.04	
0.48					
highway-mpg	-0.83	-0.72	-0.63	-0.03	
0.45					
price	0.91	0.83	0.65	0.12	
-0.18					
	horsepower	peak-rpm	city-mpg	highway-mpg	price
normalized-losses	0.24	0.30	-0.25	-0.20	0.19
wheel-base	0.50	-0.32	-0.49	-0.54	0.68
length	0.66	-0.27	-0.67	-0.70	0.81
width	0.69	-0.20	-0.69	-0.70	0.81
height	0.01	-0.30	-0.07	-0.13	0.26
curb-weight	0.81	-0.24	-0.81	-0.83	0.91
engine-size	0.82	-0.28	-0.73	-0.72	0.83
bore	0.65	-0.31	-0.62	-0.63	0.65
stroke	0.14	-0.07	-0.04	-0.03	0.12
compression-ratio	-0.36	-0.03	0.48	0.45	-0.18
horsepower	1.00	0.11	-0.91	-0.88	0.85
peak-rpm	0.11	1.00	-0.13	-0.06	-0.08
city-mpg	-0.91	-0.13	1.00	0.97	-0.83
highway-mpg	-0.88	-0.06	0.97	1.00	-0.83
price	0.85	-0.08	-0.83	-0.83	1.00

Conclusiones EDA

Análisis gráfico y descriptivo

- Dataset de tamaño 205x26.
- 26 atributos: 11 Categóricos y 15 Numéricos.

- el 59% de los registros en el dataset tiene una clasificacion de riesgo 0 y 1.
- No hay carros bastante seguros (symbolizing = -3) en el dataset.
- el pago de pérdida promedio relativo por año de vehículo asegurado toma valores entre 65 y 256 con una media de 122.
- El top 5 de marcas de vehiculos con mas registros en este dataset son Toyota, Nissan, Mazda, Honda y Mitsubishi.
- El 90,24% de los vehiculos en este dataset tienen motores a gasolina.
- El 18,05% de los vehiculos en este dataset tienen motores con turbo.
- El dataset se encuentra bien balanceado entre carros con 4 puertas (55,61%) y 2 puertas (43,41%).
- La mayoría de vehiculos en este dataset tiene carroceria tipo Sedan (46,83%).
- El 58,54% de vehiculos en este dataset tiene traccion delantera y solo es 4,39% es 4x4.
- El 98.54% de vehiculos en este dataset tiene el motor ubicado en la parte frontal.
- Los vehiculos en este dataset tienen una media de distancia entre ejes de 98,75 (posiblemente en cm).
- Los vehiculos en este dataset tienen una media de longitud de 174,04, una media de anchura de 65.90 y una media de altura de 53,72 (posiblemente cm).
- El peso medio de los vehiculos en este dataset es de 2555,56 (posiblemente kilogramos).
- El 72,20% de los vehiculos en este dataset cuenta con motor con árbol de levas en cabeza.
- El 77,56% de los vehiculos en este dataset cuenta con un motor de 4 cilindros.
- El 45,85% de los vehiculos en este dataset cuenta con un sistema multipuerto de inyección de combustible electrónica y el 32,20% tiene un sistema de combustible 2BBL (dos barriles).
- La media de caballos de fuerza de los vehiculos en este dataset es de 5125, con un maximo de 6660 y un minimo de 4150.
- La media de millas por galon de los vehiculos en este dataset es de 25,21 mpg en ciudad y de 30,75 mpg en carretera.
- Segun este dataset El carro con mejor rendimiento de millas por galon en ciudad es de marca Honda y El carro con peor rendimiento de millas por galon en ciudad es de marca Jaguar.
- Segun este dataset El carro con mejor rendimiento de millas por galon en carretera es de marca Honda y El carro con peor rendimiento de millas por galon en carretera es de marca Mercedez.
- La media de precio de los vehiculos de este data set es de 13207,12 (posiblemente dolares). Siendo el vehiculo mas costoso un Mercedez benz y el mas barato un Subaru.


Correlaciones Spearman

Las variables independientes estan altamente correlacionadas en su mayoria, aqui mencionaremos las correlaciones mas fuertes:

- wheel base y length: 0.91
- wheel base y width: 0.81
- wheel base y curb weigth: 0.77
- Length y width: 0.89


- Length y curb weight: 0.89
- Length y engine size: 0.78
- Length y price: 0.81
- Length y highway mpg: -0.70
- Width y curb weight: 0.86
- Width y engine size: 0.77
- Width y highway mpg: -0.70
- Width y price: 0.81
- curb-weight y engine size: 0.88
- curb-weight y bore: 0.70
- curb-weight y horse power: 0.81
- curb-weight y city mpg: -0.81
- curb-weight y highway mpg: -0.83
- curb-weight y price: 0.91
- engine-size y bore: 0.72
- engine-size y horse power: 0.82
- engine-size y city mpg: -0.73
- engine-size y highway mpg: -0.72
- engine-size y price: 0.83
- highway mpg y city mpg: 0.97

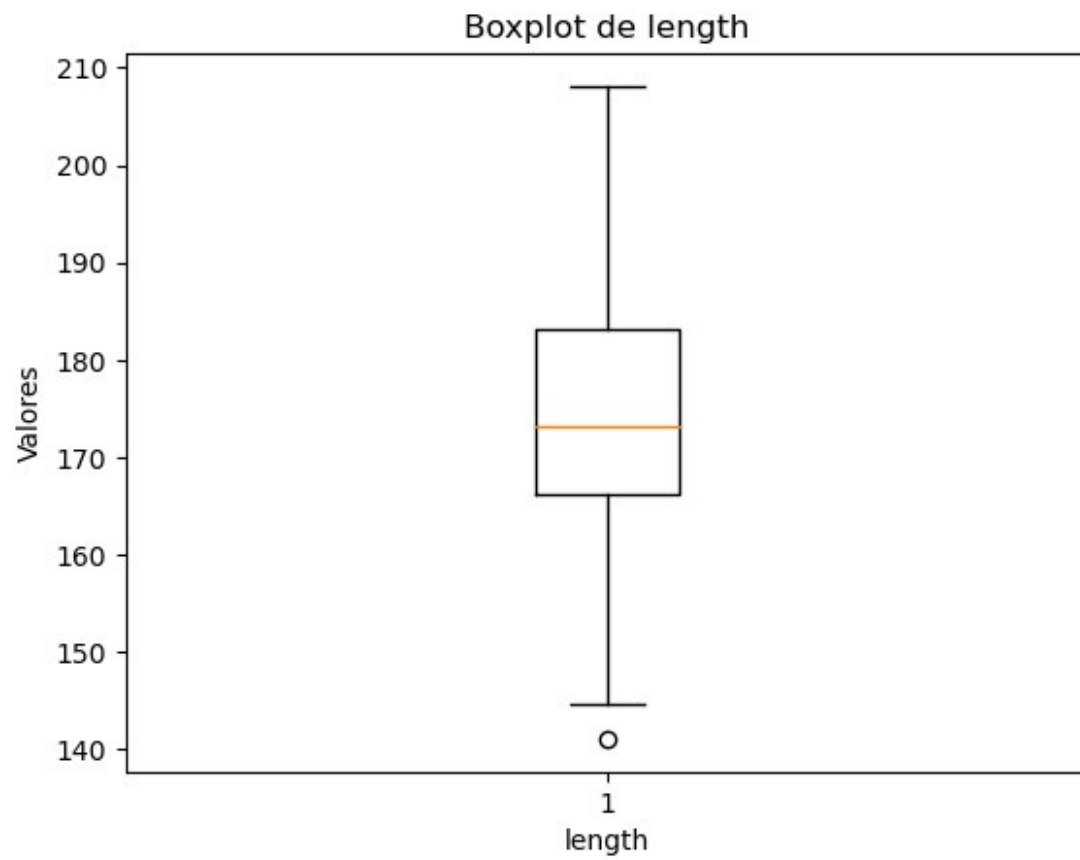
```
#Generamos graficas boxplots para las variables numericas con el fin de analizar datos atipicos
numeric_columns = data.select_dtypes(include=['int', 'float']).columns
for column in numeric_columns:
    plt.boxplot(data[column])
    plt.xlabel(column)
    plt.ylabel('Valores')
    plt.title('Boxplot de ' + column)
    plt.show()
```








Boxplot de wheel-base














Boxplot de engine-size




Boxplot de bore




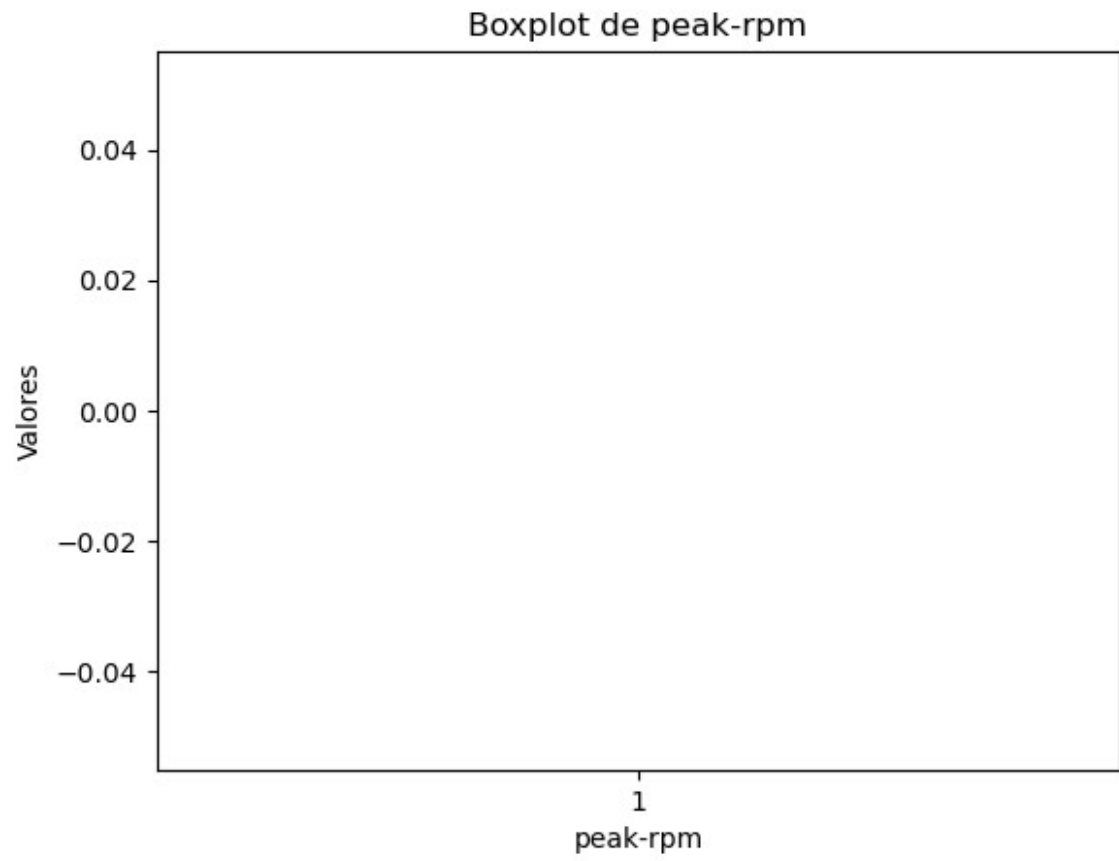
Boxplot de stroke




Boxplot de compression-ratio

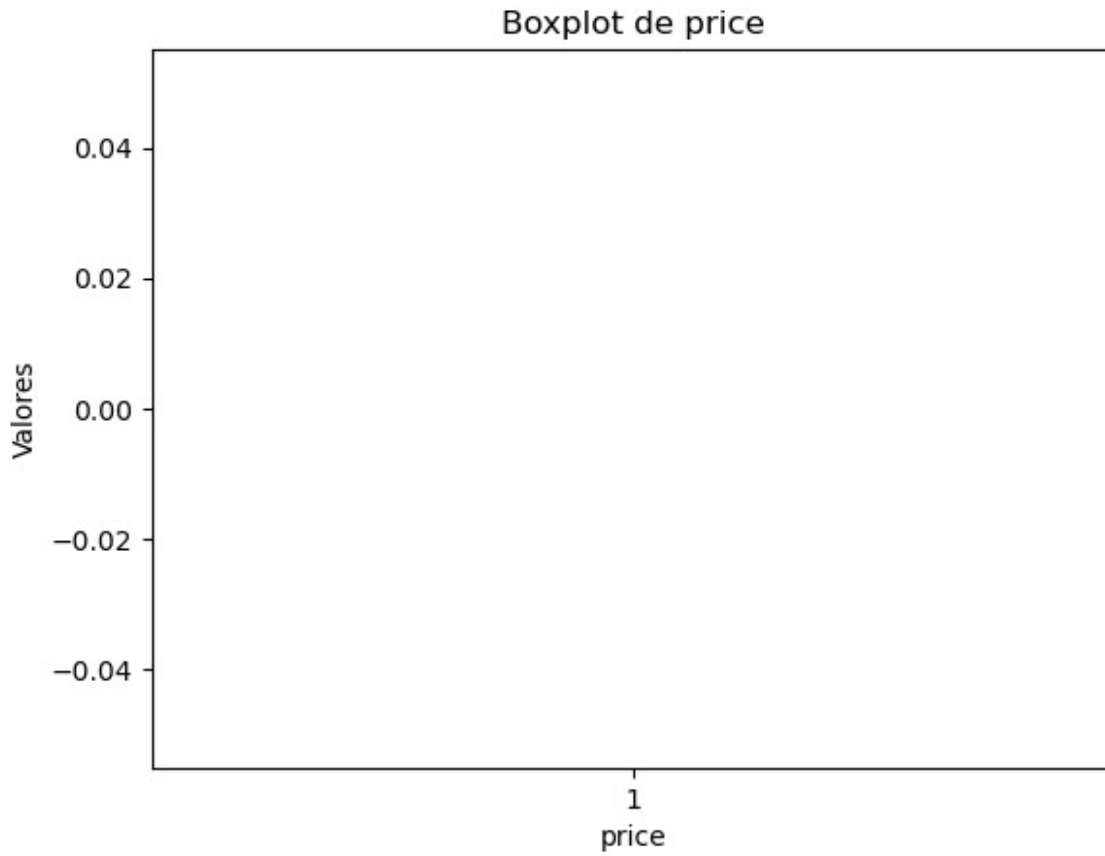












#Los atributos cuantitativos son los siguientes:

numeric_columns

```
Index(['symboling', 'normalized-losses', 'wheel-base', 'length',  
      'width',  
      'height', 'curb-weight', 'engine-size', 'bore', 'stroke',  
      'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg',  
      'highway-mpg', 'price'],  
      dtype='object')
```

*#Creamos un dataset nuevo que contiene solamente los atributos
numericos:*

```
data2=data[['normalized-losses', 'wheel-base', 'length', 'width',  
            'height',  
            'curb-weight', 'engine-size', 'bore', 'stroke', 'compression-  
ratio',  
            'horsepower', 'peak-rpm', 'city-mpg', 'highway-mpg', 'price']]
```

#Normalizamos las escalas de todas las variables cuantitativas

```
normalizada= pd.DataFrame(preprocessing.scale(data2))
```

```
normalizada.columns=['normalized-losses', 'wheel-base', 'length',  
                    'width', 'height',  
                    'curb-weight', 'engine-size', 'bore', 'stroke', 'compression-
```

```
ratio',
      'horsepower', 'peak-rpm', 'city-mpg', 'highway-mpg', 'price']
```

normalizada

	normalized-losses	wheel-base	length	width	height	curb-weight
\						
0	NaN	-1.69	-0.43	-0.84	-2.02	-0.01
1	NaN	-1.69	-0.43	-0.84	-2.02	-0.01
2	NaN	-0.71	-0.23	-0.19	-0.54	0.51
3	1.19	0.17	0.21	0.14	0.24	-0.42
4	1.19	0.11	0.21	0.23	0.24	0.52
..
200	-0.76	1.72	1.20	1.40	0.73	0.76
201	-0.76	1.72	1.20	1.35	0.73	0.95
202	-0.76	1.72	1.20	1.40	0.73	0.88
203	-0.76	1.72	1.20	1.40	0.73	1.27
204	-0.76	1.72	1.20	1.40	0.73	0.98

	engine-size	bore	stroke	compression-ratio	horsepower	peak-
rpm \						
0	0.07	0.51	-1.82	-0.29	0.17	-
0.26						
1	0.07	0.51	-1.82	-0.29	0.17	-
0.26						
2	0.60	-2.38	0.68	-0.29	1.26	-
0.26						
3	-0.43	-0.51	0.46	-0.04	-0.06	
0.78						
4	0.22	-0.51	0.46	-0.54	0.27	
0.78						
..
.						
200	0.34	1.65	-0.33	-0.16	0.25	
0.57						
201	0.34	1.65	-0.33	-0.36	1.41	
0.37						
202	1.11	0.92	-1.22	-0.34	0.75	
0.78						
203	0.44	-1.17	0.46	3.24	0.04	-

```
0.68
204      0.34  1.65  -0.33      -0.16      0.25
0.57
```

	city-mpg	highway-mpg	price
0	-0.65	-0.55	0.04
1	-0.65	-0.55	0.42
2	-0.95	-0.69	0.42
3	-0.19	-0.11	0.09
4	-1.11	-1.27	0.54
..
200	-0.34	-0.40	0.46
201	-0.95	-0.84	0.74
202	-1.11	-1.13	1.04
203	0.12	-0.55	1.17
204	-0.95	-0.84	1.19

```
[205 rows x 15 columns]
```

```
normalizada.std(axis=0)
```

normalized-losses	1.00
wheel-base	1.00
length	1.00
width	1.00
height	1.00
curb-weight	1.00
engine-size	1.00
bore	1.00
stroke	1.00
compression-ratio	1.00
horsepower	1.00
peak-rpm	1.00
city-mpg	1.00
highway-mpg	1.00
price	1.00

```
dtype: float64
```

```
normalizada.mean(axis=0)
```

normalized-losses	0.00
wheel-base	-0.00
length	0.00
width	0.00
height	-0.00
curb-weight	0.00
engine-size	0.00
bore	0.00
stroke	0.00
compression-ratio	-0.00

```
horsepower      -0.00
peak-rpm        0.00
city-mpg        0.00
highway-mpg     0.00
price           -0.00
dtype: float64
```

```
#Eliminamos la columna Normalized losses ya que consideramos no es relevante para la construccion del modelo,
#ademas tiene un 20% de missing values
```

```
normalizada = normalizada.drop("normalized-losses", axis=1)
```

```
#Eliminamos los valores Na de todas las columnas
normalizada=normalizada.dropna()
```

```
#Realizamos componentes principales con el nuevo dataset normalizado
pca = PCA()
pca.fit(normalizada)
```

```
PCA()
```

```
componentes_coef=pd.DataFrame(pca.components_)
```

```
pca.explained_variance_
```

```
array([7.51333306, 2.26460554, 1.20343145, 0.8950368 , 0.60919366,
        0.41239635, 0.31705757, 0.27127324, 0.1208348 , 0.10924733,
        0.08305822, 0.06249769, 0.05215879, 0.01924072])
```

```
var_exp=pca.explained_variance_ratio_ # varianza explicada por cada PC
cum_var_exp = np.cumsum(var_exp) # varianza acumulada por los primeros n PCs
```

```
var_exp
```

```
array([0.5392332 , 0.16253113, 0.08637048, 0.06423694, 0.04372193,
        0.02959776, 0.02275528, 0.01946933, 0.00867233, 0.0078407 ,
        0.0059611 , 0.00448547, 0.00374345, 0.00138091])
```


```
cum_var_exp
```

```
array([0.5392332 , 0.70176432, 0.7881348 , 0.85237175, 0.89609368,
        0.92569144, 0.94844671, 0.96791604, 0.97658837, 0.98442907,
        0.99039018, 0.99487565, 0.99861909, 1.          ])
```

```
dataPca = pca.transform(normalizada)
```

```
plt.figure(figsize=(15, 7))
plt.bar(range(len(var_exp)), var_exp, alpha=0.3333, align='center',
label='Varianza explicada por cada PC', color = 'blue')
plt.step(range(len(cum_var_exp)), cum_var_exp,
where='mid',label='Varianza explicada acumulada')
plt.ylabel('Porcentaje de varianza explicada')
```

```
plt.xlabel('Componentes principales')
plt.legend(loc='best')
plt.show()
```



```
np.sum(pca.explained_variance_ratio_[0:4])
0.8523717464905413
```

Los primeros 6 componentes principales juntos logran explicar el 85,23% de la varianza original

```
principalComponents = pca.fit_transform(normalizada)

#Creamos un dataset con todos los Eigenvectores
component_columns = ['principal component {}'.format(i) for i in
range(1, 15)]
principalDf = pd.DataFrame(data=principalComponents,
columns=component_columns)

principalDf
```

	principal component 1	principal component 2	principal component
3 \			
0	-0.66	-2.14	
0.25			
1	-0.53	-2.17	
0.20			
2	0.39	-1.32	-
1.46			
3	-0.18	-0.23	-
0.07			

4	1.25	-1.17	-
0.05			
..
..			
190	2.62	0.39	
1.09			
191	3.43	-0.25	
1.03			
192	3.45	-0.52	
1.44			
193	2.39	3.04	-
1.12			
194	3.26	0.08	
1.08			
	principal component 4	principal component 5	principal component
6 \			
0	2.42	0.17	
0.09			
1	2.47	0.31	
0.14			
2	-0.64	0.41	
1.94			
3	-1.13	0.32	-
0.14			
4	-1.18	0.12	
0.29			
..
..			
190	-0.45	0.09	-
0.97			
191	-0.33	-0.00	-
0.83			
192	-0.31	0.88	-
0.06			
193	-0.53	2.01	
0.60			
194	-0.45	0.23	-
0.85			
	principal component 7	principal component 8	principal component
9 \			
0	-0.52	-0.97	-
0.17			
1	-0.38	-0.87	-
0.08			
2	0.65	-1.39	-
0.18			
3	-0.17	-0.08	

0.16		
4	0.13	-0.72
0.23		
..
..		
190	-0.85	1.05
0.06		
191	-0.49	0.59
0.01		
192	-0.53	0.55
0.06		-
193	-0.58	-1.03
0.09		-
194	-0.65	0.72
0.30		

	principal component 10	principal component 11	principal component 12 \
--	------------------------	------------------------	--------------------------

0	-0.14	0.45
0.36		
1	-0.30	0.63
0.30		
2	0.45	-0.05
-0.38		
3	-0.10	0.30
-0.22		
4	-0.41	0.07
0.18		
..
...		
190	-0.05	-0.19
-0.02		
191	0.23	-0.28
-0.78		
192	-0.26	-0.37
0.08		
193	-0.41	-0.05
-0.28		
194	-0.56	0.14
-0.08		

	principal component 13	principal component 14
--	------------------------	------------------------

0	0.15	0.01
1	0.20	0.04
2	0.08	-0.13
3	-0.26	0.02
4	0.10	0.07
..
190	0.03	-0.04

```

191             0.01             -0.07
192            -0.35             -0.12
193            -0.00              0.27
194             0.14             -0.16

[195 rows x 14 columns]


finalDf = pd.concat([principalDf, data[['symboling']]], axis = 1)
data["symboling"].unique()
array([ 3,  1,  2,  0, -1, -2])

def biplot(data, loadings, index1, index2, labels=None):
    plt.figure(figsize=(15, 7))
    xs = data[:,index1]
    ys = data[:,index2]
    n=loadings.shape[0]
    scalex = 1.0/(xs.max()- xs.min())
    scaley = 1.0/(ys.max()- ys.min())
    plt.scatter(xs*scalex,ys*scaley)
    for i in range(n):
        plt.arrow(0, 0, loadings[i,index1],
loadings[i,index2],color='r',alpha=0.5)
        if labels is None:
            plt.text(loadings[i,index1]* 1.15, loadings[i,index2] *
1.15, "Var"+str(i+1), color='g', ha='center', va='center')
        else:
            plt.text(loadings[i,index1]* 1.15, loadings[i,index2] *
1.15, labels[i], color='g', ha='center', va='center')
    plt.xlim(-1,1)
    plt.ylim(-1,1)
    plt.xlabel("PC{}".format(index1))
    plt.ylabel("PC{}".format(index2))
    plt.grid()

normalizada.columns
Index(['wheel-base', 'length', 'width', 'height', 'curb-weight',
'engine-size',
      'bore', 'stroke', 'compression-ratio', 'horsepower', 'peak-
rpm',
      'city-mpg', 'highway-mpg', 'price'],
      dtype='object')

# PC0 VS PC1
biplot(dataPca, pca.components_, 0, 1,['wheel-base', 'length',
'width', 'height', 'curb-weight', 'engine-size',
      'bore', 'stroke', 'compression-ratio', 'horsepower', 'peak-
rpm',
      'city-mpg', 'highway-mpg', 'price'])


```



En el biplot se presenta el círculo unitario con la representación de variables y observaciones, se evidencia que la variable mejor representada en el componente PC0 es stroke. Las demás variables parecen ser explicadas por una combinación lineal de PC0 y PC1.

Las variables width, price, engine-size, curb-weight y highway-mpg parecen tener poca significancia debido a la corta longitud de su línea.

```
# PC0 VS PC2
biplot(dataPca, pca.components[, 0, 2], ['wheel-base', 'length',
'width', 'height', 'curb-weight', 'engine-size',
'bore', 'stroke', 'compression-ratio', 'horsepower', 'peak-
rpm',
'city-mpg', 'highway-mpg', 'price'])
```



En el biplot se presenta el círculo unitario con la representación de variables y observaciones, se evidencia que la variable mejor representada por PC0 es city mpg. Las demás variables parecen ser explicadas por una combinación lineal de PC0 Y PC2.


Las variables price, curb-weight, width y highway-mpg parecen tener poca significancia debido a la corta longitud de su línea.

```
normalizada.columns
```

```
Index(['wheel-base', 'length', 'width', 'height', 'curb-weight',  
      'engine-size',  
      'bore', 'stroke', 'compression-ratio', 'horsepower', 'peak-  
rpm',  
      'city-mpg', 'highway-mpg', 'price'],  
      dtype='object')
```

```
#PC0 vs PC3
```

```
biplot(dataPca, pca.components_, 0, 3, ['wheel-base', 'length',  
    'width', 'height', 'curb-weight', 'engine-size',  
    'bore', 'stroke', 'compression-ratio', 'horsepower', 'peak-  
rpm',  
    'city-mpg', 'highway-mpg', 'price'])
```



En el biplot se presenta el círculo unitario con la representación de variables y observaciones, se evidencia que la variable mejor representada por PC0 es stroke y city mpg. Las demás variables parecen ser explicadas por una combinación lineal de PC0 Y PC3.

Las variables que se decide eliminar por que no contribuyen en la creación de las componentes: price, curb-weight, highway-mpg, width, engine-size

Modelo

```
data=pd.read_csv("imports-85.data", header=None)
data.columns = [
    "symboling", "normalized-losses", "make", "fuel-type",
    "aspiration", "num-of-doors", "body-style", "drive-wheels",
    "engine-location", "wheel-base", "length", "width", "height",
    "curb-weight", "engine-type", "num-of-cylinders",
    "engine-size", "fuel-system", "bore", "stroke", "compression-
ratio", "horsepower", "peak-rpm", "city-mpg", "highway-mpg",
    "price"
]
```

```
data=data.replace("?",np.nan)
```

```
data.drop("normalized-losses", axis=1)
```

	symboling	make	fuel-type	aspiration	num-of-doors	body-style
0	3	alfa-romero	gas	std	two	convertible
1	3	alfa-romero	gas	std	two	convertible
2	1	alfa-romero	gas	std	two	hatchback
3	2	audi	gas	std	four	sedan
4	2	audi	gas	std	four	sedan
..
200	-1	volvo	gas	std	four	sedan
201	-1	volvo	gas	turbo	four	sedan
202	-1	volvo	gas	std	four	sedan
203	-1	volvo	diesel	turbo	four	sedan
204	-1	volvo	gas	turbo	four	sedan

	drive-wheels	engine-location	wheel-base	length	...	engine-size
0	rwd	front	88.60	168.80	...	130
1	rwd	front	88.60	168.80	...	130
2	rwd	front	94.50	171.20	...	152
3	fwd	front	99.80	176.60	...	109

4	4wd	front	99.40	176.60	...	136
..
200	rwd	front	109.10	188.80	...	141
201	rwd	front	109.10	188.80	...	141
202	rwd	front	109.10	188.80	...	173
203	rwd	front	109.10	188.80	...	145
204	rwd	front	109.10	188.80	...	141
	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm
city-mpg \						
0	mpfi	3.47	2.68	9.00	111	5000
21						
1	mpfi	3.47	2.68	9.00	111	5000
21						
2	mpfi	2.68	3.47	9.00	154	5000
19						
3	mpfi	3.19	3.40	10.00	102	5500
24						
4	mpfi	3.19	3.40	8.00	115	5500
18						
..
...						
200	mpfi	3.78	3.15	9.50	114	5400
23						
201	mpfi	3.78	3.15	8.70	160	5300
19						
202	mpfi	3.58	2.87	8.80	134	5500
18						
203	idi	3.01	3.40	23.00	106	4800
26						
204	mpfi	3.78	3.15	9.50	114	5400
19						
	highway-mpg	price				
0	27	13495				
1	27	16500				
2	26	16500				
3	30	13950				
4	22	17450				
..				
200	28	16845				
201	25	19045				

```
202      23  21485
203      27  22470
204      25  22625
```

```
[205 rows x 25 columns]
```

```
data["price"]=data["price"].astype(float)
data["bore"]=data["bore"].astype(float)
data["stroke"]=data["stroke"].astype(float)
data["horsepower"]=data["horsepower"].astype(float)
data["peak-rpm"]=data["peak-rpm"].astype(float)
data["price"]=data["price"].astype(float)
data["symboling"]=data["symboling"].astype(object)
```

```
data["symboling"].value_counts()
```

```
0      67
1      54
2      32
3      27
-1     22
-2       3
```

```
Name: symboling, dtype: int64
```

```
#Agrupamos los vehiculos segun su riesgo para reducir el numero de categorias
```

```
data['Target'] = np.where((data['symboling'] >= 1) &
                          (data['symboling'] <= 3), 'Riesgoso',
                          np.where((data['symboling'] >= -3) &
                          (data['symboling'] <= 0), 'Seguro', 'Otro'))
```

```
data["Target"].value_counts()
```

```
Riesgoso    113
```

```
Seguro       92
```

```
Name: Target, dtype: int64
```

```
## al tener una distribucion balanceada, se comienza a modelar
#price, curb-weight, highway-mpg, width, engine-size
#make, num-of-doors, body-style (agrupar sedan vs otras),fuel-system
(idi y las otras)
```

```
#Agrupamos los vehiculos con body style hardtop y convertible juntos
```

```
data["body-style"] = data["body-style"].replace({'sedan': 'sedan',
'hatchback': 'hatchback', 'wagon':
'wagon', 'hardtop': 'others', 'convertible': 'others'})
```

```
#Agrupamos los vehiculos con fuel system idi, lbbl, spdi, 4bbl, mfi y spfi juntos
```

```
data["fuel-system"] = data["fuel-system"].replace({'mpfi': 'mpfi',
'2bbl': '2bbl', 'idi':
```

```
'others', 'lbbbl': 'others', 'spdi': 'others', '4bbl': 'others', 'mfi': 'others', 'spfi': 'others'})
```

```
#Aplicamos one hot encoding para las variables categoricas seleccionadas
```

```
encoded_data = pd.get_dummies(data[["body-style", "make", "num-of-doors", "fuel-system"]])
```

```
#Se concatenan ambos datasets
```

```
datas=pd.concat([data[['wheel-base', 'length', 'height', 'bore', 'stroke', 'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg']], encoded_data], axis=1)
```

```
# Se eliminan los valores na
```

```
datas=datas.dropna()
```

```
#Definimos las variables independientes y la variable dependiente
```

```
X=datas
```

```
y=data["Target"]
```

```
X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 199 entries, 0 to 204
```

```
Data columns (total 40 columns):
```

#	Column	Non-Null Count	Dtype
0	wheel-base	199 non-null	float64
1	length	199 non-null	float64
2	height	199 non-null	float64
3	bore	199 non-null	float64
4	stroke	199 non-null	float64
5	compression-ratio	199 non-null	float64
6	horsepower	199 non-null	float64
7	peak-rpm	199 non-null	float64
8	city-mpg	199 non-null	int64
9	body-style_hatchback	199 non-null	uint8
10	body-style_others	199 non-null	uint8
11	body-style_sedan	199 non-null	uint8
12	body-style_wagon	199 non-null	uint8
13	make_alfa-romero	199 non-null	uint8
14	make_audi	199 non-null	uint8
15	make_bmw	199 non-null	uint8
16	make_chevrolet	199 non-null	uint8
17	make_dodge	199 non-null	uint8
18	make_honda	199 non-null	uint8
19	make_isuzu	199 non-null	uint8
20	make_jaguar	199 non-null	uint8
21	make_mazda	199 non-null	uint8
22	make_mercedes-benz	199 non-null	uint8
23	make_mercury	199 non-null	uint8


```

24 make_mitsubishi      199 non-null    uint8
25 make_nissan           199 non-null    uint8
26 make_peugot          199 non-null    uint8
27 make_plymouth        199 non-null    uint8
28 make_porsche          199 non-null    uint8
29 make_renault          199 non-null    uint8
30 make_saab             199 non-null    uint8
31 make_subaru           199 non-null    uint8
32 make_toyota           199 non-null    uint8
33 make_volkswagen       199 non-null    uint8
34 make_volvo            199 non-null    uint8
35 num-of-doors_four     199 non-null    uint8
36 num-of-doors_two      199 non-null    uint8
37 fuel-system_2bbl      199 non-null    uint8
38 fuel-system_mpf       199 non-null    uint8
39 fuel-system_others    199 non-null    uint8
dtypes: float64(8), int64(1), uint8(31)
memory usage: 21.6 KB

```

#Esta libreria nos permite correr varios modelos ed Machine learning al tiempo y comparar sus metricas.

```

X, y = make_classification(n_samples=1000, n_features=20,
random_state=42)

```

```

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

```

```

clf = LazyClassifier(verbose=0, ignore_warnings=True)
models, predictions = clf.fit(X_train, X_test, y_train, y_test)
models

```

```

100%|██████████| 29/29 [00:01<00:00, 22.03it/s]

```

	Accuracy	Balanced Accuracy	ROC AUC
F1 Score \			
Model			
RandomForestClassifier	0.90	0.90	0.90
0.90			
XGBClassifier	0.89	0.89	0.89
0.89			
BaggingClassifier	0.89	0.89	0.89
0.88			
LGBMClassifier	0.89	0.89	0.89
0.89			
DecisionTreeClassifier	0.88	0.88	0.88
0.88			
ExtraTreesClassifier	0.87	0.87	0.87
0.87			

AdaBoostClassifier	0.87	0.87	0.87
0.87			
CalibratedClassifierCV	0.86	0.87	0.87
0.87			
LinearSVC	0.86	0.86	0.86
0.86			
LinearDiscriminantAnalysis	0.85	0.86	0.86
0.85			
RidgeClassifierCV	0.85	0.86	0.86
0.85			
RidgeClassifier	0.85	0.86	0.86
0.85			
LogisticRegression	0.85	0.86	0.86
0.86			
NuSVC	0.84	0.85	0.85
0.84			
SVC	0.84	0.85	0.85
0.85			
Perceptron	0.83	0.83	0.83
0.83			
NearestCentroid	0.82	0.83	0.83
0.82			
BernoulliNB	0.81	0.81	0.81
0.80			
SGDClassifier	0.81	0.81	0.81
0.81			
GaussianNB	0.80	0.80	0.80
0.79			
PassiveAggressiveClassifier	0.80	0.80	0.80
0.80			
ExtraTreeClassifier	0.79	0.80	0.80
0.79			
QuadraticDiscriminantAnalysis	0.80	0.79	0.79
0.79			
KNeighborsClassifier	0.79	0.79	0.79
0.78			
LabelSpreading	0.77	0.78	0.78
0.77			
LabelPropagation	0.77	0.78	0.78
0.77			
DummyClassifier	0.47	0.50	0.50
0.30			
Time Taken			
Model			
RandomForestClassifier	0.19		
XGBClassifier	0.12		
BaggingClassifier	0.09		
LGBMClassifier	0.11		

DecisionTreeClassifier	0.02
ExtraTreesClassifier	0.12
AdaBoostClassifier	0.19
CalibratedClassifierCV	0.11
LinearSVC	0.03
LinearDiscriminantAnalysis	0.02
RidgeClassifierCV	0.01
RidgeClassifier	0.02
LogisticRegression	0.01
NuSVC	0.04
SVC	0.03
Perceptron	0.01
NearestCentroid	0.01
BernoulliNB	0.01
SGDClassifier	0.01
GaussianNB	0.01
PassiveAggressiveClassifier	0.02
ExtraTreeClassifier	0.01
QuadraticDiscriminantAnalysis	0.01
KNeighborsClassifier	0.02
LabelSpreading	0.05
LabelPropagation	0.04
DummyClassifier	0.01

Random Forest parece ser el mejor modelo, vamos a comprobarlo

MODELO 1 RANDOM FOREST

```
forest = RandomForestClassifier(criterion='gini',
                               n_estimators=5,
                               random_state=1,
                               n_jobs=2)

forest.fit(X_train, y_train)

y_pred = forest.predict(X_test)
print('Accuracy: %.3f' % accuracy_score(y_test, y_pred))

Accuracy: 0.850

from sklearn.metrics import cohen_kappa_score
kappa = cohen_kappa_score(y_test, y_pred)
print("Coeficiente de Kappa:", kappa)

Coeficiente de Kappa: 0.7010463378176383
```

Utilizando un modelo Random Forest obtenemos una clasificacion con un accuracy de 85% y un Kappa de 0.701, lo que nos permite concluir que tenemos un buen modelo mucho mejor que el baseline

```

feature_importance = forest.feature_importances_

#Obtenemos la relevancia de las variables predictoras
feature_importance

array([0.02159787, 0.1497592 , 0.03012585, 0.01857064, 0.00620078,
        0.18341466, 0.01649461, 0.01845453, 0.02005582, 0.01282518,
        0.02649634, 0.03763293, 0.02080572, 0.01594713, 0.07845535,
        0.01681059, 0.02133975, 0.01126815, 0.28096426, 0.01278065])

X

array([[ -0.6693561 , -1.49577819, -0.87076638, ..., -1.26733697,
        -1.2763343 ,  1.01664321],
       [ 0.09337237,  0.78584826,  0.10575379, ..., -0.12270893,
        0.6934308 ,  0.91136272],
       [-0.90579721, -0.60834121,  0.29514098, ...,  0.83049813,
        -0.73733198, -0.5782121 ],
       ...,
       [-0.20013455, -1.46108168,  1.79701652, ..., -1.50280171,
        -1.27473745,  1.60111869],
       [ 0.03935575,  0.24868361, -0.47532342, ...,  0.09912579,
        0.54269228,  1.20827474],
       [ 0.76921528,  0.47076539,  0.16994471, ...,  0.65611162 ,
        0.64333186, -2.02100232]])

import plotly.graph_objects as go
import plotly.io as pio

pio.renderers.default = "notebook" # Elige el renderizador por defecto (cambia a "notebook" si usas Jupyter Notebook)

fig = go.Figure(data=go.Bar(x=datas.columns, y=feature_importance))
fig.update_layout(title="Feature Importance Random Forest",
xaxis_title="Feature", yaxis_title="Importance")
fig.show()

```

Podemos concluir que las variables predictoras mas relevantes son: Si el vehiculo es de marca honda, compression ratio, Longitud del vehiculo y si el vehiculo es de marca Audi

MODELO 2 XGBOOST

```

!pip install xgboost

Requirement already satisfied: xgboost in c:\users\mqa200-0489\
anaconda3\lib\site-packages (1.7.5)
Requirement already satisfied: scipy in c:\users\mqa200-0489\
anaconda3\lib\site-packages (from xgboost) (1.9.1)
Requirement already satisfied: numpy in c:\users\mqa200-0489\
anaconda3\lib\site-packages (from xgboost) (1.21.5)

```

```

from xgboost import XGBClassifier

modelGB = XGBClassifier(gamma=2.1784702961406848,
                        learning_rate=0.06271852908557515,
                        max_depth=4,
                        n_estimators=45,
                        subsample=0.7856261494444738)
modelGB.fit(X_train, y_train)

XGBClassifier(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None,
              feature_types=None,
              gamma=2.1784702961406848, gpu_id=None, grow_policy=None,
              importance_type=None, interaction_constraints=None,
              learning_rate=0.06271852908557515, max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=4, max_leaves=None,
              min_child_weight=None, missing=nan,
              monotone_constraints=None,
              n_estimators=45, n_jobs=None, num_parallel_tree=None,
              predictor=None, random_state=None, ...)

# make predictions for test data
y_pred = modelGB.predict(X_test)
predictions = [round(value) for value in y_pred]

# evaluate predictions
accuracy = accuracy_score(y_test, predictions)
print("Accuracy: %.2f%%" % (accuracy * 100.0))

Accuracy: 90.50%

kappa = cohen_kappa_score(y_test, predictions)
print("Coeficiente de Kappa:", kappa)

Coeficiente de Kappa: 0.8105305145592341

```

Utilizando un modelo XGBOOST una clasificacion con un accuracy de 90.5% y un Kappa de 0.810, lo que nos permite concluir que tenemos un buen modelo mucho mejor que el baseline

```

#Obtenemos la relevancia de las variables predictoras
feature_importancel = modelGB.feature_importances_

fig = go.Figure(data=go.Bar(x=datas.columns, y=feature_importancel))
fig.update_layout(title="Feature Importance XGB00ST ",
                  xaxis_title="Feature", yaxis_title="Importance")
fig.show()

```

Podemos concluir que las variables predictoras mas relevantes son: Compression ratio y Si el vehiculo es de marca honda

MODELO 3 Gradient boosting

```
gb_clf = GradientBoostingClassifier()
gb_clf.fit(X_train, y_train)

GradientBoostingClassifier()

feature_importance2 =gb_clf.feature_importances_
y_pred = gb_clf.predict(X_test)
print('Accuracy: %.3f' % accuracy_score(y_test, y_pred))

Accuracy: 0.910

kappa = cohen_kappa_score(y_test, y_pred)
print("Coeficiente de Kappa:", kappa)

Coeficiente de Kappa: 0.8206278026905829
```

Utilizando un modelo Gradient Boosting una clasificacion con un accuracy de 91% y un Kappa de 0.820, lo que nos permite concluir que tenemos un buen modelo mucho mejor que el baseline

```
fig = go.Figure(data=go.Bar(x=datas.columns, y=feature_importance2))
fig.update_layout(title="Feature Importance Gradient Boosting",
xaxis_title="Feature", yaxis_title="Importance")
fig.show()
```

Podemos concluir que las variables predictoras mas relevantes son: Compression ratio y Si el vehiculo es de marca audi

De los 3 modelos realizados, Gradient Boosting obtuvo los mejores resultados en Accuracy y kappa

Clusters

```
from sklearn.cluster import KMeans, AgglomerativeClustering
from scipy.cluster.hierarchy import dendrogram, linkage, fcluster
```

```

datas=data[['wheel-base', 'length', 'height', 'bore', 'stroke',
            'compression-ratio',
            'horsepower', 'peak-rpm', 'city-mpg']]

```

#Normalizamos las escalas de las variables

```

dataStd = pd.DataFrame(preprocessing.scale(datas))
dataStd.columns=datas.columns

```

```

dataStd.describe()

```

	wheel-base	length	height	bore	stroke	compression-ratio	\
count	205.00	205.00	205.00	201.00	201.00	205.00	
mean	-0.00	0.00	-0.00	0.00	0.00	-0.00	
std	1.00	1.00	1.00	1.00	1.00	1.00	
min	-2.02	-2.68	-2.43	-2.89	-3.75	-0.79	
25%	-0.71	-0.63	-0.71	-0.66	-0.46	-0.39	
50%	-0.29	-0.07	0.15	-0.07	0.11	-0.29	
75%	0.61	0.74	0.73	0.95	0.49	-0.19	
max	3.69	2.77	2.49	2.24	2.89	3.24	

	horsepower	peak-rpm	city-mpg
count	203.00	203.00	205.00
mean	-0.00	0.00	0.00
std	1.00	1.00	1.00
min	-1.42	-2.04	-1.87
25%	-0.86	-0.68	-0.95
50%	-0.23	0.16	-0.19
75%	0.30	0.78	0.73
max	4.64	3.08	3.64

```

mean_values = dataStd.mean(axis=0)
mean_values_formatted = mean_values.round(12)
print(mean_values_formatted)

```

```

wheel-base      -0.00
length           0.00
height          -0.00
bore             0.00
stroke           0.00
compression-ratio -0.00
horsepower       -0.00
peak-rpm         0.00
city-mpg         0.00
dtype: float64

```

```

dataStd.mean(axis=0)

```

```

wheel-base      -0.00
length           0.00
height          -0.00
bore             0.00

```

```
stroke          0.00
compression-ratio -0.00
horsepower      -0.00
peak-rpm        0.00
city-mpg        0.00
dtype: float64
```

```
dataStd.std(axis=0)
```

```
wheel-base      1.00
length          1.00
height          1.00
bore            1.00
stroke          1.00
compression-ratio 1.00
horsepower      1.00
peak-rpm        1.00
city-mpg        1.00
dtype: float64
```

```
#Eliminamos los valores na de todas las columnas
dataStd=dataStd.dropna()
```

Optmización del numeros de clusters

KMEANS

Metodo del codo:

```
WSSs = []
for i in range(1,15) :
    km = KMeans(n_clusters=i, random_state=0)
    km.fit(dataStd)
    WSSs.append(km.inertia_)
WSSs
```


```
[1785.5661430309424,
1288.275276600507,
1040.7090214947486,
896.0345069319069,
783.7834631140512,
716.8104966874055,
636.8703227885688,
581.1104240782172,
533.4665968409798,
492.2208676385021,
443.51001421691086,
423.2454229037552,
```



```
400.159891009055,  
377.8037356390405]
```

```
plt.plot(range(1, 15), WSSs)
```

```
[<matplotlib.lines.Line2D at 0x1e2763b37c0>]
```



En el criterio del codo se ve que el numero optimo de k es 2.

Calinski-Harabaz

```
CHs = []  
for i in range(2,15) :  
    km = KMeans(n_clusters=i, random_state=0)  
    km.fit(dataStd)  
    CH = calinski_harabasz_score(dataStd, km.labels_)  
    CHs.append(CH)
```


CHs

```
[76.04453991021924,  
70.14064104652849,  
64.52826972525425,  
61.98964671049615,  
57.55212589590608,  
57.71703427286113,  
56.55454325960451,
```

```
55.743629307077256,  
55.17899091427457,  
56.888580669937355,  
54.718730525831546,  
53.663041671143745,  
53.02632044508594]
```

```
plt.plot(range(2, 15), CHs)
```

```
[<matplotlib.lines.Line2D at 0x1e276b39490>]
```



Segun Calinski el K optimo es aproximadamente 3

```
def grafico(dataStd, k):  
    kmeans = KMeans(n_clusters=k, random_state=0, n_init=10)  
    kmeans.fit(dataStd)  
    y_clusters = kmeans.labels_  
    cluster_labels = np.unique(y_clusters)  
  
    silueta_puntos = silhouette_samples(dataStd, y_clusters,  
metric='euclidean')  
  
    fig, ax = plt.subplots()  
    y_ax_lower, y_ax_upper = 0, 0  
    yticks = []
```

```

colores = ['r', 'g', 'b', 'y', 'o', "w"]
for i, c in enumerate(cluster_labels):
    silueta_puntos_c = silueta_puntos[y_clusters == c]
    silueta_puntos_c.sort()
    y_ax_upper += len(silueta_puntos_c)
    color = colores[i]
    ax.barh(range(y_ax_lower, y_ax_upper), silueta_puntos_c,
height=1.0,
            edgecolor='none', color=color)

    yticks.append((y_ax_lower + y_ax_upper) / 2.)
    y_ax_lower += len(silueta_puntos_c)


silueta_promedio = np.mean(silueta_puntos)
ax.axvline(silueta_promedio, color="black", linestyle="--")

ax.set_yticks(yticks)
ax.set_yticklabels(cluster_labels + 1)
ax.set_ylabel('Cluster')
ax.set_xlabel('Coeficiente de silueta')


plt.tight_layout()
plt.show()

grafico(dataStd, 2)


```



```
grafico(dataStd, 3)
```



```
grafico(dataStd, 4)
```



Segun la silueta podemos concluir que el K optimo es 2

```
var_num = data[['wheel-base', 'length', 'height', 'bore', 'stroke',
                'compression-ratio',
                'horsepower', 'peak-rpm', 'city-mpg']]
```

```
dataStd.head()
```

	wheel-base	length	height	bore	stroke	compression-ratio
horsepower \						
0	-1.69	-0.43	-2.02	0.51	-1.82	-0.29
0.17						
1	-1.69	-0.43	-2.02	0.51	-1.82	-0.29
0.17						
2	-0.71	-0.23	-0.54	-2.38	0.68	-0.29
1.26						
3	0.17	0.21	0.24	-0.51	0.46	-0.04
0.06						
4	0.11	0.21	0.24	-0.51	0.46	-0.54
0.27						
peak-rpm						
0	-0.26					-0.65

1	-0.26	-0.65
2	-0.26	-0.95
3	0.78	-0.19
4	0.78	-1.11

#Utilizamos Cluster jerarquico con fusiones y KMeans con K = 2

```
fig, axes = plt.subplots(2,2,figsize=(15,15))
```




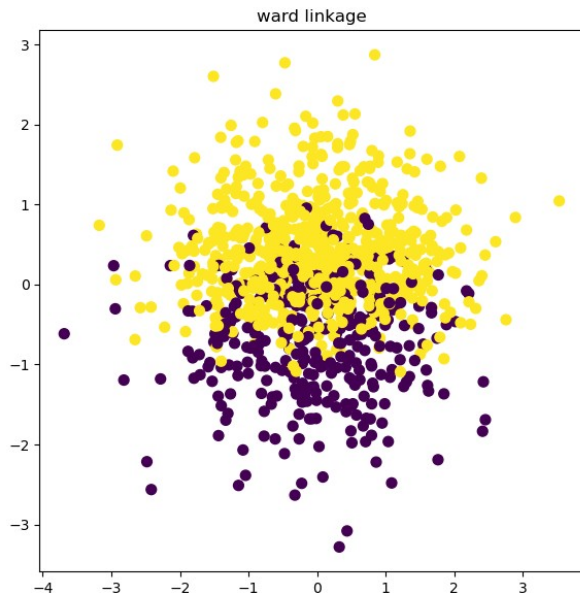
```
link = 'ward'
clustering = AgglomerativeClustering(linkage=link, n_clusters=2)
clustering.fit(X)
axes[0][0].scatter(X[:, 0], X[:, 1], c=clustering.labels_, s=50,
cmap='viridis')
axes[0][0].set_title("%s linkage" % link)
```

```
link = 'complete'
clustering = AgglomerativeClustering(linkage=link, n_clusters=2)
clustering.fit(X)
axes[0][1].scatter(X[:, 0], X[:, 1], c=clustering.labels_, s=50,
cmap='viridis')
axes[0][1].set_title("%s linkage" % link)
```

```
link = 'average'
clustering = AgglomerativeClustering(linkage=link, n_clusters=2)
clustering.fit(X)
axes[1][0].scatter(X[:, 0], X[:, 1], c=clustering.labels_, s=50,
cmap='viridis')
axes[1][0].set_title("%s linkage" % link)
```

```
clustering = KMeans(n_clusters=2)
clustering.fit(X)
axes[1][1].scatter(X[:, 0], X[:, 1], c=clustering.labels_, s=50,
cmap='viridis')
axes[1][1].set_title("K-Means")
```

```
Text(0.5, 1.0, 'K-Means')
```



Analizando las graficas podemos concluir que con $K=2$, Ward y Kmeans obtuvieron los mejores resultados identificando los clusters

Dendrograma


```
#Dendrograma
plt.figure(figsize=(16, 7))
plt.title('Dendrograma de clustering jerárquico')
plt.xlabel('Registro')
plt.ylabel('Distancia')
dendrogram(fusiones,
            orientation='top',
```



```

distance_sort='descending',
truncate_mode='level',
p=3,
show_leaf_counts=True)
plt.show()

```



Observando el Dendrograma se puede evidenciar que con K=2 los clusters quedan bien balanceados

#Obtenemos el cluster de cada punto con Clustering jerárquico

```

k = 2
clusters = fcluster(fusiones, k, criterion='maxclust')
clusters

```

```

array([2, 1, 2, 1, 1, 1, 2, 1, 2, 1, 1, 1, 1, 1, 2, 1, 1, 1, 2, 1,
1,
1, 1, 2, 1, 1, 1, 1, 2, 2, 1, 1, 2, 1, 2, 1, 1, 1, 1, 2, 1, 2,
1,
1, 2, 1, 1, 1, 1, 2, 1, 2, 1, 1, 2, 1, 2, 2, 1, 1, 1, 1, 1, 1,
2,
1, 2, 1, 1, 2, 1, 2, 1, 1, 1, 1, 1, 2, 2, 2, 1, 1, 1, 1, 1, 2,
1,
1, 1, 2, 2, 1, 2, 1, 1, 1, 2, 2, 1, 2, 2, 1, 2, 1, 2, 2, 1, 1,
1,
1, 2, 2, 2, 1, 1, 1, 2, 1, 2, 1, 2, 1, 1, 1, 2, 2, 1, 2, 1, 1,
2,
1, 2, 1, 1, 1, 1, 2, 2, 1, 2, 1, 1, 1, 2, 1, 1, 1, 2, 2, 1, 1,
1,
2, 1, 1, 2, 2, 2, 2, 1, 1, 2, 1, 2, 1, 1, 2, 1, 2, 1, 1, 1, 1,
2,
1, 1, 1, 2, 2, 1, 2, 1, 2, 1, 2, 1, 1, 1, 2, 1, 2, 1, 1, 1, 1,

```

2,	1, 1, 1, 1, 2, 1, 2, 1, 1, 2, 2, 2, 2, 1, 2, 1, 1, 2, 2, 1, 2,
2,	2, 1, 2, 1, 1, 1, 1, 1, 2, 2, 2, 2, 2, 2, 1, 1, 2, 2, 2, 2, 2,
1,	1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 1, 1, 1, 1, 2, 1, 1, 1, 1, 2,
1,	1, 2, 2, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 2, 1, 1, 2, 1, 2, 1,
2,	1, 1, 1, 2, 1, 2, 1, 1, 1, 1, 1, 1, 2, 2, 2, 1, 1, 2, 1, 1, 2,
1,	1, 2, 1, 2, 1, 1, 1, 1, 2, 2, 2, 1, 1, 1, 2, 2, 1, 1, 2, 1, 1,
1,	1, 1, 2, 2, 2, 2, 1, 1, 2, 1, 1, 1, 1, 1, 2, 2, 2, 2, 2, 1, 1,
2,	2, 2, 2, 1, 1, 1, 1, 2, 1, 1, 2, 2, 1, 1, 1, 1, 1, 2, 2, 1, 1,
1,	1, 2, 1, 1, 2, 2, 1, 1, 2, 1, 1, 2, 1, 1, 1, 1, 1, 1, 1, 2, 1,
1,	1, 1, 2, 1, 1, 1, 2, 1, 1, 1, 1, 2, 2, 2, 1, 2, 2, 1, 1, 2, 2,
2,	1, 2, 1, 2, 1, 1, 1, 2, 1, 1, 1, 1, 1, 2, 2, 1, 1, 2, 1, 2, 2,
1,	1, 1, 1, 2, 2, 1, 2, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 2, 1, 2,
1,	1, 2, 2, 1, 1, 1, 2, 1, 1, 2, 2, 1, 2, 1, 1, 1, 1, 1, 1, 1, 1,
2,	2, 2, 1, 1, 1, 1, 2, 2, 2, 1, 2, 2, 1, 1, 1, 1, 1, 1, 2, 1, 1,
1,	1, 1, 1, 2, 1, 1, 2, 2, 2, 1, 1, 1, 1, 2, 1, 1, 1, 2, 2, 1, 2,
1,	1, 2, 2, 1, 2, 1, 1, 1, 1, 1, 1, 2, 1, 2, 1, 1, 2, 2, 1, 2, 2,
1,	1, 1, 1, 1, 2, 1, 2, 2, 1, 2, 1, 2, 2, 2, 1, 2, 1, 1, 2, 1, 1,
1,	1, 1, 1, 1, 1, 1, 2, 1, 2, 1, 1, 1, 1, 1, 1, 2, 1, 2, 1, 1, 1,
2,	1, 1, 1, 2, 2, 1, 2, 1, 2, 1, 2, 1, 1, 1, 2, 1, 1, 1, 1, 2, 2,
1,	1, 1, 1, 1, 1, 1, 1, 2, 1, 2, 2, 1, 1, 1, 1, 1, 2, 2, 1, 2, 1,
1,	2, 1, 1, 1, 1, 1, 1, 2, 1, 1, 2, 1, 1, 1, 1, 1, 2, 2, 1, 1, 2,
1,	2, 1, 2, 1, 2, 1, 1, 1, 1, 1, 2, 1, 1, 1, 2, 2, 1, 1, 1, 1, 1,
1,	1, 2, 1, 2, 1, 1, 1, 2, 1, 2, 1, 1, 1, 2, 1, 2, 1, 1, 1, 1, 2,
2,	2, 2, 1, 1, 2, 1, 1, 1, 1, 1, 1, 2, 2, 1, 2, 2, 1, 2, 1, 1,
1,	

```

1,      2, 1, 1, 1, 2, 1, 1, 2, 1, 2, 1, 1, 1, 2, 1, 2, 1, 2, 1, 1, 2,
1,      2, 2, 1, 2, 1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 1, 1, 1, 1, 2, 2, 1,
1,      1, 1, 1, 1, 1, 1, 2, 1, 1, 1, 1, 1, 1, 1, 2, 1, 2, 1, 2, 1, 2,
2,      1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 2, 2, 1, 1, 1, 1, 1,
2,      1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2,
2,      1, 2, 2, 2, 1, 1, 2, 2, 1, 1, 1, 1, 2, 2, 1, 1, 1, 1, 1, 1, 2,
2,      1, 2, 2, 1, 2, 1, 2, 1, 1, 1, 1, 2, 2, 1, 1, 1, 2, 1, 1, 2, 2,
1,      2, 2, 1, 1, 1, 2, 2, 1, 1, 1, 1, 2, 1, 1, 2, 1, 1, 1, 2, 1, 2,
1,      1, 1, 1, 2, 1, 2, 2, 1, 1, 1, 2, 2, 2, 1, 1, 1, 1, 2, 1, 2, 1,
2,      2, 1, 1, 1, 1, 2, 2, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 1, 1, 1,
2,      1, 2, 1, 2, 2, 2, 2, 1, 1, 1, 1, 1, 2, 1, 1, 1, 1, 1, 1, 2, 2,
2,      1, 1, 2, 1, 1, 1, 1, 1, 2, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 1,
1,      1, 2, 2, 1, 1, 1, 1, 2, 1, 1], dtype=int32)

```

#Clustering usando KMEANS con K=2

```
kmeans = KMeans(n_clusters=2, random_state=0, n_init=10)
```

```
kmeans.fit(dataStd)
```

```
KMeans(n_clusters=2, random_state=0)
```

```
kmeans.labels_
```

```

array([1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1,
1,
1,      1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1,      1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0,
0,      0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 1,
1,      1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,      0, 0, 0, 0, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
0,      1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1,
1,      1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0,
0,
0,

```

```

0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0])

clusters = kmeans.predict(dataStd)
clusters

array([1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1,
1,
1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1,
1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0,
0,
0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 1, 1,
1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
0, 0, 0, 0, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
0,
1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1,
1,
1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0,
0,
1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
0])

counter=Counter(clusters)
print(counter)

Counter({1: 101, 0: 98})

```

Tenemos 2 Clusters, El primero con 101 registros y el segundo con 98 registros

```

dataStd.loc[:, 'Cluster'] = clusters


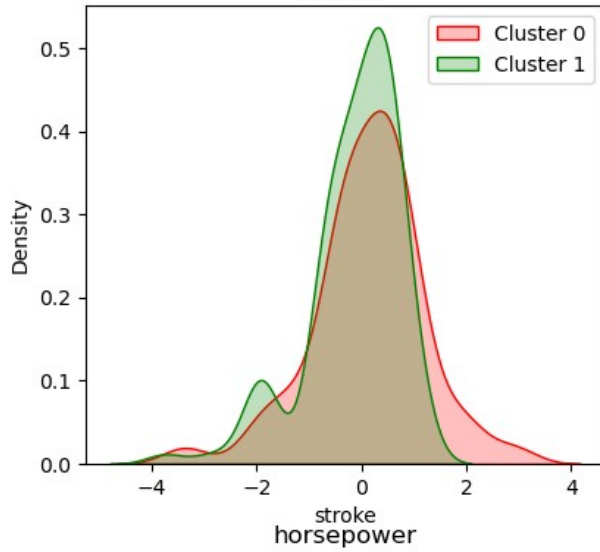
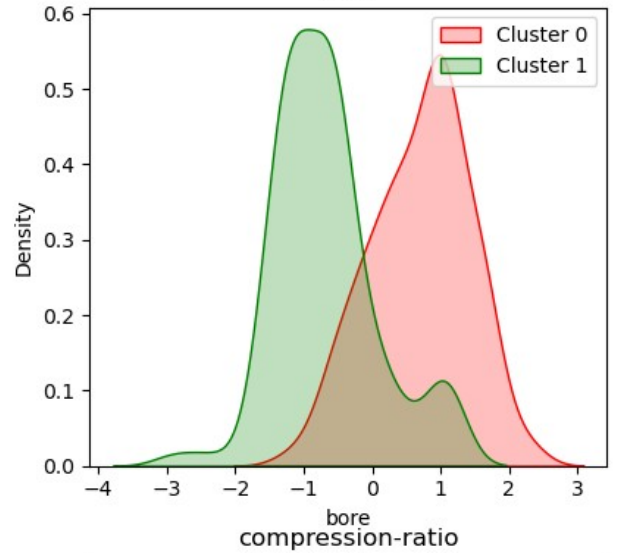



#Valores promedio de cada cluster
df_agrupado = Nuevo.groupby('label').median()
df_agrupado

```

	wheel-base	length	height	bore	stroke	compression-ratio	\
label							
0	94.49	166.28	53.00	3.15	3.23	9.10	
1	102.41	183.52	55.00	3.56	3.35	8.90	

	horsepower	peak-rpm	city-mpg
label			
0	69.92	5200.18	30.01
1	116.03	4999.69	18.98

```
fig = plt.figure(figsize=(10,500))
i=1
for var in var_num:
    ax = fig.add_subplot(math.ceil(len(var_num)/2), 2, i)
    sns.kdeplot(dataStd.loc[dataStd.Cluster==0][var], shade=True,
color='r', ax=ax);
    sns.kdeplot(dataStd.loc[dataStd.Cluster==1][var], shade=True,
color='g', ax=ax);
    sns.kdeplot(dataStd.loc[dataStd.Cluster==2][var], shade=True,
color='b', ax=ax);
    plt.title(var)
    plt.legend(['Cluster 0', 'Cluster 1', 'Cluster 2'])
    i+=1
```



```

fig = plt.figure(figsize=(15,15))
colorPalette = ["r", "g"]
ax = fig.add_subplot(2, 2, 1)
sns.scatterplot(x="compression-ratio", y="length", hue="Cluster",
data=dataStd, ax=ax, palette=colorPalette, s=100, alpha=0.5)
plt.title("compression-ratio vs. length")
ax = fig.add_subplot(2, 2, 2)
sns.scatterplot(x="compression-ratio", y="horsepower", hue="Cluster",
data=dataStd, ax=ax, palette=colorPalette, s=100, alpha=0.5)
plt.title("compression-ratio vs. horsepower")

Text(0.5, 1.0, 'compression-ratio vs. horsepower')

```

